

Hadoop 1.X

Instructor: Jayant Mohite

Results: _____

Name: _____

Class: Wisdom Sprouts

Period: Hadoop

Date: 10th December 2015

1. How many blocks will be created for a file of size 300 MB provided the default block size is 128 MB and rest all default configurations are set.
 - a. 15
 - b. 3
 - c. 9
 - d. 6

2. What is the role of a Name Node
 - a. Manage Storage at Master Level
 - b. Manage Processing at Master Level
 - c. Manage Storage at Slave Level
 - d. Manage Processing at Slave Level

3. What is the role of a Data Node
 - a. Manage Storage at Master Level
 - b. Manage Processing at Master Level
 - c. Manage Storage at Slave Level
 - d. Manage Processing at Slave Level

4. What is the role of a Job Tracker
 - a. Manage Storage at Master Level
 - b. Manage Processing at Master Level

- c. Manage Storage at Slave Level
 - d. Manage Processing at Slave Level
5. What is the role of a Task Tracker
- a. Manage Storage at Master Level
 - b. Manage Processing at Master Level
 - c. Manage Storage at Slave Level
 - d. Manage Processing at Slave Level
6. If you have to change the replication factor value, then you will modify this property in
- a. hdfs-site.xml
 - b. core-site.xml
 - c. mapred-site.xml
 - d. hadoop-env.sh
7. If you have to change the number of Reducers, then you will modify this property in
- a. hdfs-site.xml
 - b. core-site.xml
 - c. mapred-site.xml
 - d. hadoop-env.sh
8. If you have to change the master server, then you will modify this property in
- a. hdfs-site.xml
 - b. core-site.xml
 - c. mapred-site.xml
 - d. hadoop-env.sh
9. The default Replication factor enforces, presence of ____ copies of every block
- a. 1
 - b. 2

- c. 3
- d. 4

10. Shuffle and Sort phase is internal to

- a. Map Phase
- b. Reduce Phase
- c. HDFS
- d. Linux

11. The command "for service in /etc/init.d/hadoop-hdfs ; do sudo \$service start;" will initiate how many services

- a. 2
- b. 3
- c. 4
- d. 5

12. The command "for service in /etc/init.d/hadoop-0.20-mapreduce ; do sudo \$service start;" will initiate how many services

- a. 2
- b. 3
- c. 4
- d. 5

13. The command "for service in /etc/init.d/hadoop* ; do sudo \$service start;" will initiate how many services

- a. 2
- b. 3
- c. 4
- d. 5

14. If Data Node 1 contains blocks "a" and "b" and if Data Node 2 contains block "c", then incase Data Node 2 goes down, then Task Tracker 1 can process block "c"
- a. True
 - b. False
15. If Name Node goes down, then the Secondary Name Node takes its place
- a. True
 - b. False
16. If the storage capacity of a server is 10 TB, then the HDFS Capacity of that server is also
- a. 10 TB
 - b. 7 TB
 - c. 3 TB
 - d. Can't define